

PUFFERFISH : COMMUNICATION- EFFICIENT MODELS AT NO EXTRA COST

MLSys 2021

**DML GROUP MEETING
12.10**

OUTLINE

- Introduction
- Effective deep factorized network training
- Strategies for mitigating training accuracy
- Experiments

Introduction

- 通信开销导致数据并行的分布式训练无法达到最佳的加速效果
- 通信开销归因于频繁的梯度更新
- 模型更高的准确率也导致了参数规模的扩大，加剧了通信开销
- 现有解决办法：低精度训练、梯度稀疏化
- 梯度压缩存在的问题：
 - 1) 梯度压缩的计算代价很大，例如对每个批次的梯度做SVD矩阵分解
 - 2) 现有的梯度压缩方法要么没有充分利用，要么需要额外内存
 - 3) 在现有的深度学习框架上整合这些梯度压缩方法需要大量的工作

Introduction

- 能否将梯度压缩步骤的纳入模型架构本身?
- 通信效率可以在没有额外成本的情况下得以提升
- 基于此想法设计了Pufferfish
- 设计思想
 - 1) 将模型架构分解，训练分解后的模型，一劳永逸
 - 2) 对分解后的模型采用额外的训练策略提高准确率

Effective deep factorized network training

Low-rank factorization for FC layers

- 两层的FC network可表示为：

$$h(x) = \sigma(W_1 \sigma(W_2 x))$$

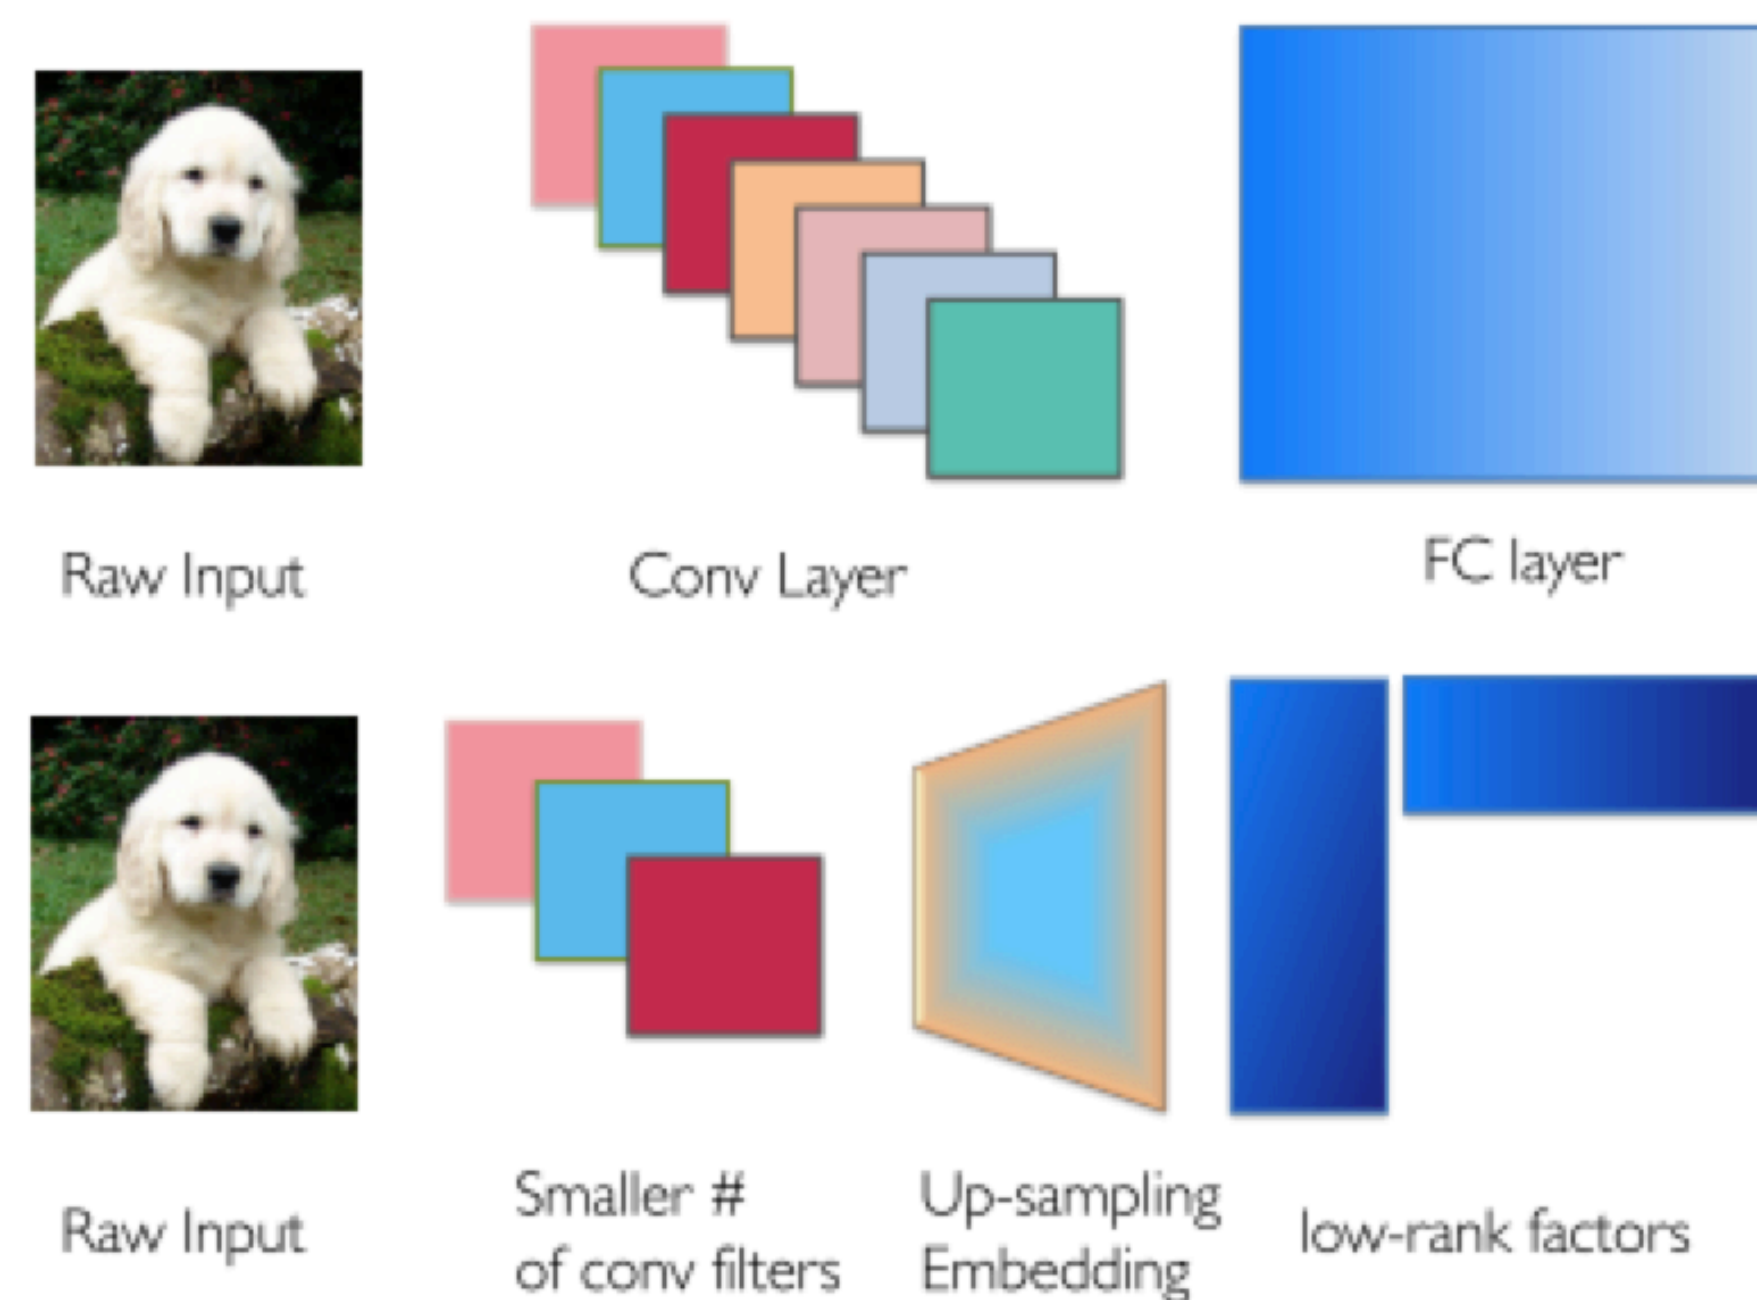
- 将每层的参数分解为

$$W_l \xrightarrow{\text{分解}} U_l V_l^T$$

Effective deep factorized network training

Low-rank factorization for convolution layers

- 原始参数形式: $W \in \mathbb{R}^{c_{in} \times c_{out} \times k \times k}$
- 整理: $W_{unrolled} \in \mathbb{R}^{c_{in} k^2 \times c_{out}}$
- 分解: $U \in \mathbb{R}^{c_{in} k^2 \times r}, V^T \in \mathbb{R}^{r \times c_{out}}$
- 转换: $U \in \mathbb{R}^{c_{in} \times r \times k \times k}, V_l^T \in \mathbb{R}^{r \times c_{out} \times 1 \times 1}$



Effective deep factorized network training

Low-rank factorization for LSTM layers

$$\begin{aligned}i_t &= \sigma(W_{ii}x_t + b_{ii} + W_{hi}h_{t-1} + b_{hi}) \\f_t &= \sigma(W_{if}x_t + b_{if} + W_{hf}h_{t-1} + b_{hf}) \\g_t &= \tanh(W_{ig}x_t + b_{ig} + W_{hg}h_{t-1} + b_{hg}) \\o_t &= \sigma(W_{io}x_t + b_{io} + W_{ho}h_{t-1} + b_{ho}) \\c_t &= f_t \odot c_{t-1} + i_t \odot g_t \\h_t &= o_t \odot \tanh(c_t).\end{aligned}$$

(1)

分解



$$\begin{aligned}i_t &= \sigma(U_{ii}V_{ii}^\top x_t + b_{ii} + U_{hi}V_{hi}^\top h_{t-1} + b_{hi}) \\f_t &= \sigma(U_{if}V_{if}^\top x_t + b_{if} + U_{hf}V_{hf}^\top h_{t-1} + b_{hf}) \\g_t &= \tanh(U_{ig}V_{ig}^\top x_t + b_{ig} + U_{hg}V_{hg}^\top h_{t-1} + b_{hg}) \\o_t &= \sigma(U_{io}V_{io}^\top x_t + b_{io} + U_{ho}V_{ho}^\top h_{t-1} + b_{ho}) \\c_t &= f_t \odot c_{t-1} + i_t \odot g_t \\h_t &= o_t \odot \tanh(c_t).\end{aligned} \quad (2)$$

Effective deep factorized network training

Low-rank network factorization for Transformer

- Multi-Head Attention: $\text{MultiHead}(Q, K, V) = \text{Concat}(\text{head}_1, \dots, \text{head}_p)W^O$
where $\text{head}_i = \text{Attention}(QW_i^Q, KW_i^K, VW_i^V)$.
- FFN: $\text{FFN}(x) = \max(0, xW_1 + b_1)W_2 + b_2$
- 可分解参数: $W_i^Q, W_i^K, W_i^V \in \mathbb{R}^{pd \times d}$ $W^O \in \mathbb{R}^{pd \times pd}$ $W_1 \in \mathbb{R}^{pd \times 4pd}, W_2 \in \mathbb{R}^{4pd \times pd}$


Effective deep factorized network training

Computational complexity and model size

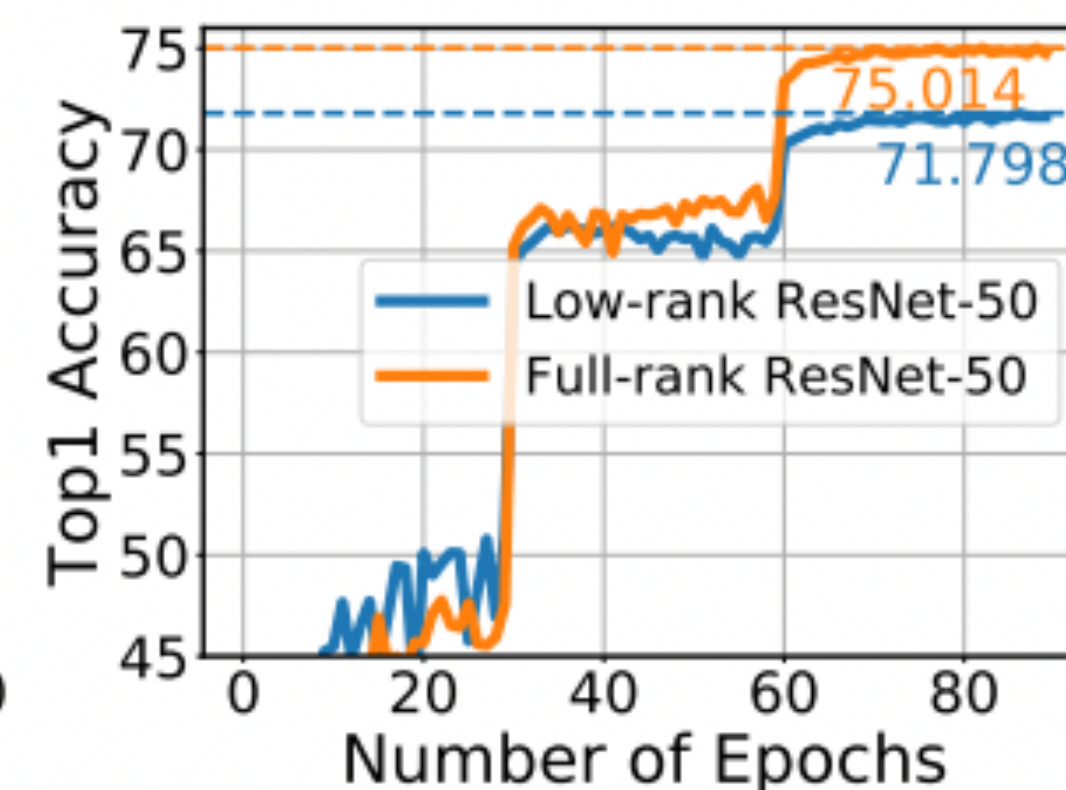
Networks	# Params.	Computational Complexity
Vanilla FC	$m \times n$	$\mathcal{O}(mn)$
Factorized FC	$r(m + n)$	$\mathcal{O}(r(m + n))$
Vanilla Conv.	$c_{\text{in}} \times c_{\text{out}} \times k^2$	$\mathcal{O}(c_{\text{in}} c_{\text{out}} k^2 HW)$
Factorized Conv.	$c_{\text{in}} r k^2 + r c_{\text{out}}$	$\mathcal{O}(r c_{\text{in}} k^2 HW + r HW c_{\text{out}})$
Vanilla LSTM	$4(dh + h^2)$	$\mathcal{O}(dh + h^2)$
Factorized LSTM	$4dr + 12hr$	$\mathcal{O}(dr + hr)$
Vanilla Attention	$4p^2 d^2$	$\mathcal{O}(Np^2 d^2 + N^2 d)$
Factorized Attention	$(3p + 5)prd$	$\mathcal{O}(rpdN + N^2 d)$
Vanilla FFN	$8p^2 d^2$	$\mathcal{O}(p^2 d^2 N)$
Factorized FFN	$10pdr$	$\mathcal{O}(rpdN)$

Strategies for mitigating training accuracy

- 模型分解一定程度上导致准确率下降
- 采取训练策略尽可能减少对准确率的影响
 - 1) 混合网络架构 (hybrid network architecture)
 - 2) vanilla warm-up training




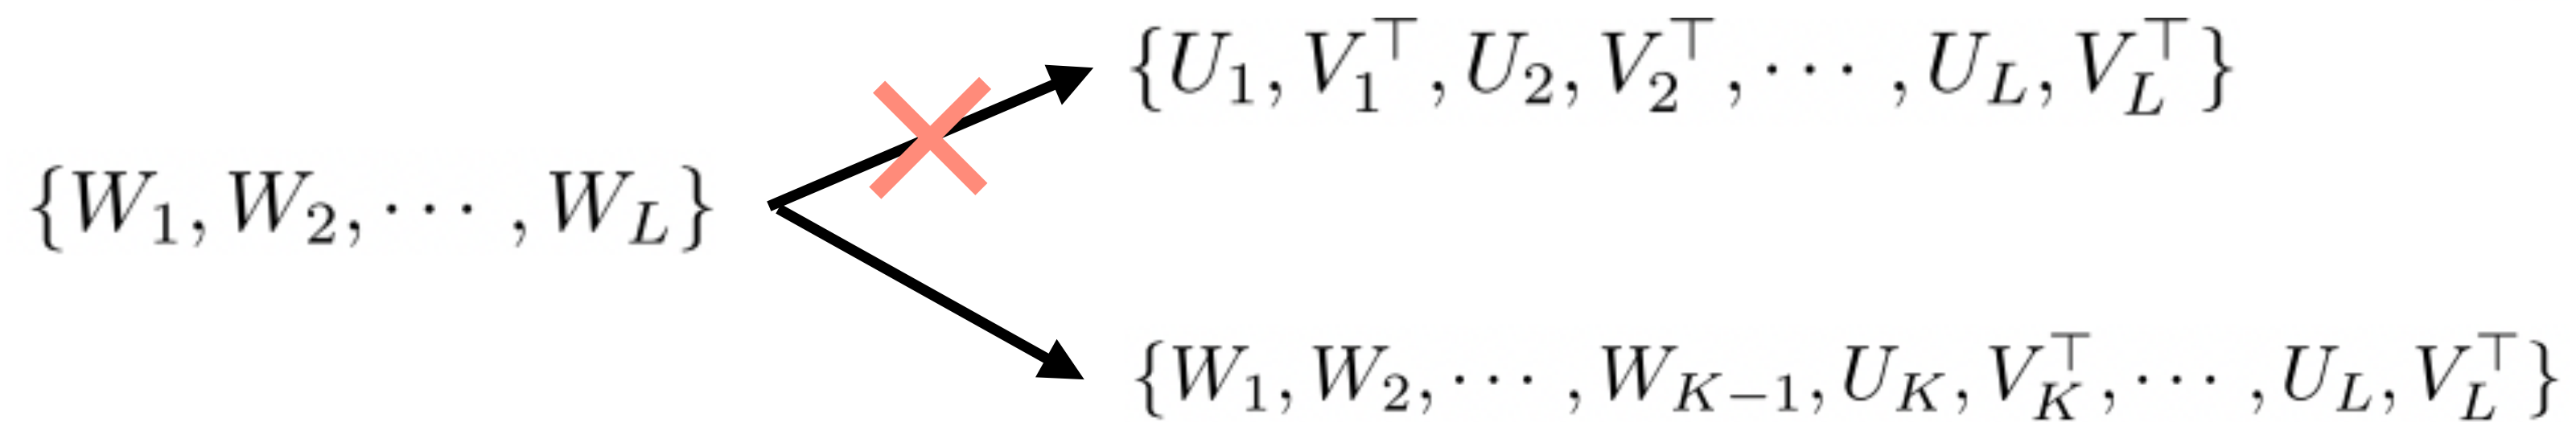
(a) VGG-11 on CIFAR-10



(b) ResNet-50 on ImageNet

hybrid network architecture


- 矩阵分解不可避免地带来了误差
- 误差会随着前向传播逐渐累积
- 深度神经网络中，后层的参数数量决定整个网络的规模
- 解决方案：仅对后层layer都做分解
- 注意：最后一层FC layer通常和最后的输出结果有关，最好不要分解该层
- 缺点：引入超参数K，需要对所有类型的layer调整K



(a) Hybrid network

vanilla warm-up training

- 训练早期的epoch对最终模型的准确率至关重要
- 直接训练分解后的模型会导致精度损失，并且这种损失不能在后续训练阶段减轻
- 解决方案：使用部分训练的原始满秩模型来初始化分解后的低秩模型
- 缺点：引入超参数Warmup Epochs



(b) Vanilla warm-up training

Algorithm

Algorithm 1 PUFFERFISH Training Procedure

Input : Randomly initialized weights of vanilla N -layer architectures $\{W_1, W_2, \dots, W_L\}$, and the associated weights of hybrid N -layer architecture $\{W_1, W_2, \dots, W_{K-1}, U_K, V_K^\top, \dots, U_L, V_L^\top\}$, the entire training epochs E , the vanilla warm-up training epochs E_{wu} , and learning rate schedule $\{\eta_t\}_{t=1}^E$

Output : Trained hybrid L -layer architecture weights $\{\hat{W}_1, \hat{W}_2, \dots, \hat{W}_{K-1}, \hat{U}_K, \hat{V}_K^\top, \dots, \hat{U}_L, \hat{V}_L^\top\}$

```
for  $t \in \{1, \dots, E_{wu}\}$  do
    Train  $\{W_1, W_2, \dots, W_L\}$  with learning rate schedule  $\{\eta_t\}_{t=1}^{E_{wu}}$ ; // vanilla warm-up training
end
```

```
for  $l \in \{1, \dots, L\}$  do
    if  $l < K$  then
        copy the partially trained  $W_l$  weight to the hybrid network;
    else
         $\tilde{U}_l \Sigma_l \tilde{V}_l^\top = \text{SVD}(W_l)$ ; // Decomposing the vanilla warm-up trained weights
         $U_l = \tilde{U}_l \Sigma_l^{\frac{1}{2}}, V_l^\top = \Sigma_l^{\frac{1}{2}} \tilde{V}_l^\top$ 
    end
end
for  $t \in \{E_{wu} + 1, \dots, E\}$  do
    Train the hybrid network weights, i.e.,  $\{W_1, W_2, \dots, W_{K-1}, U_K, V_K^\top, \dots, U_L, V_L^\top\}$  with learning rate schedule  $\{\eta_t\}_{t=E_{wu}}^E$ ; // consecutive low rank training
end
```

Experiments

- Pufferfish可以训练得到一个比其他方法小**3.35倍**的模型
- 与POWERSGD、SIGNUM和Vanilla SGD相比，在CIFAR-10上训练的ResNet-18的端到端加速比分别为**1.22倍**、**1.52倍**和**1.74倍**，同时达到与Vanilla SGD相同的精度。
- 在ImageNet数据集上，与EB Train相比，Pufferfish模型的参数减少了**1.3M**，而Top-1测试精度提高了**1.76%**
- 与LTH相比，在CIFAR-10上为VGG-19实现相同的模型压缩率但加快了**5.67倍**

Conclusion

- 缺点：
 - 1) 引入超参数， r 、 K 、Warmup Epochs，部分超参数需要按layer类型调整
 - 2) 每种模型如何进行模型分解需要提前分析，在完全了解模型后才能设计出好的分解方案

THANKS!